

MLOps Assignment 2 Final Report

Student Name: Niketh Varma Tirumalaraju

Roll Number: G25AIT2142

Course: MLOps | PGD AI Program | IIT Jodhpur

Project: Hugging Face Fine-Tuning, Experiment Tracking & Model Deployment

1. Model Selection Rationale

In this project, DistilBERT was selected as the pre-trained transformer model for text classification. DistilBERT is a

lighter and faster version of BERT that maintains strong performance while reducing computational requirements. It is

well-suited for Kaggle GPU environments and provides efficient fine-tuning for NLP tasks such as genre classification.

The model was loaded using Hugging Face Transformers and fine-tuned on the Goodreads reviews dataset.

2. Kaggle Training Setup

The training workflow was implemented using a Kaggle Notebook with GPU acceleration enabled. Kaggle Secrets

were used to securely store WANDB_API_KEY and HF_TOKEN credentials. Internet access was enabled to

download Hugging Face models and upload trained artifacts. Weights & Biases (W&B;) was integrated for experiment

tracking and monitoring.

3. Training Workflow

The dataset was tokenized using DistilBertTokenizerFast with truncation and padding applied. A custom PyTorch

dataset class was used for training and evaluation. The Hugging Face Trainer API was configured with

TrainingArguments including logging, evaluation, and checkpoint saving. W&B; automatically tracked training loss,

validation loss, accuracy, and F1-score throughout training.

4. Evaluation Results

After training, the model was evaluated on the test dataset. Metrics including evaluation loss, accuracy, and weighted

F1-score were recorded. A classification report was generated and uploaded to W&B; as an artifact for reproducibility

and experiment tracking.

5. Results Table

Replace the placeholder values below with your actual evaluation metrics from W&B; or `trainer.evaluate()`.

Metric Score

Accuracy 0.58437

F1 Score 0.58258

Eval Loss 1.17489

6. Challenges & Learnings

During the implementation process, package compatibility issues and GPU memory limitations were encountered.

These were resolved by installing compatible versions of transformers and huggingface_hub and reducing the training

batch size when necessary. The project provided practical understanding of the complete MLOps workflow including

experiment tracking, reproducibility, model deployment, and cloud-based GPU training.

7. Project Links

Project Links GitHub Repository: <https://github.com/niketh532002/MLOPS-ASSIGN2>

Kaggle Notebook: <https://www.kaggle.com/code/nikethvarma/notebook12a2641f7a>

Weights & Biases Dashboard: <https://wandb.ai/nikethv6-iitj/mlops-assignment2/runs/jek60kbf?nw=nwusernikethv6>

Hugging Face Model: <https://huggingface.co/Niketh0503/distilbert-goodreads-genres>